ELIZA · answers without and with retrieval on the val split (KAN-99)

Val questions: 30 (27 with gold provisions); the test split is not used

Retrieval (bge-m3 0.8 + BM25 0.2, 8 chunks): Recall@8 0.59, Recall@5 0.43 (W&B run v7g5hhxh)

Without retrieval: Claude Haiku 5.5 (claude.ai chat), 12/11/2026, whole consolidated text of 27/07/2026 in the chat

With retrieval: Llama-3.1-8B-Instruct Q4_K_M (Ollama 0.40.15), 8 chunks of bge-m3 0.8 + BM25 0.2, temperature 0

Expected provision cited (lenient): Claude Haiku 5.5, whole text 0.81, Llama-3.1-8B + 8 chunks 0.70

Cites a provision: Claude Haiku 5.5, whole text 0.97, Llama-3.1-8B + 8 chunks 0.77; with retrieval, cites only its 8 chunks: 0.83

Says the text does not answer: Claude Haiku 5.5, whole text 0.03, Llama-3.1-8B + 8 chunks 0.33

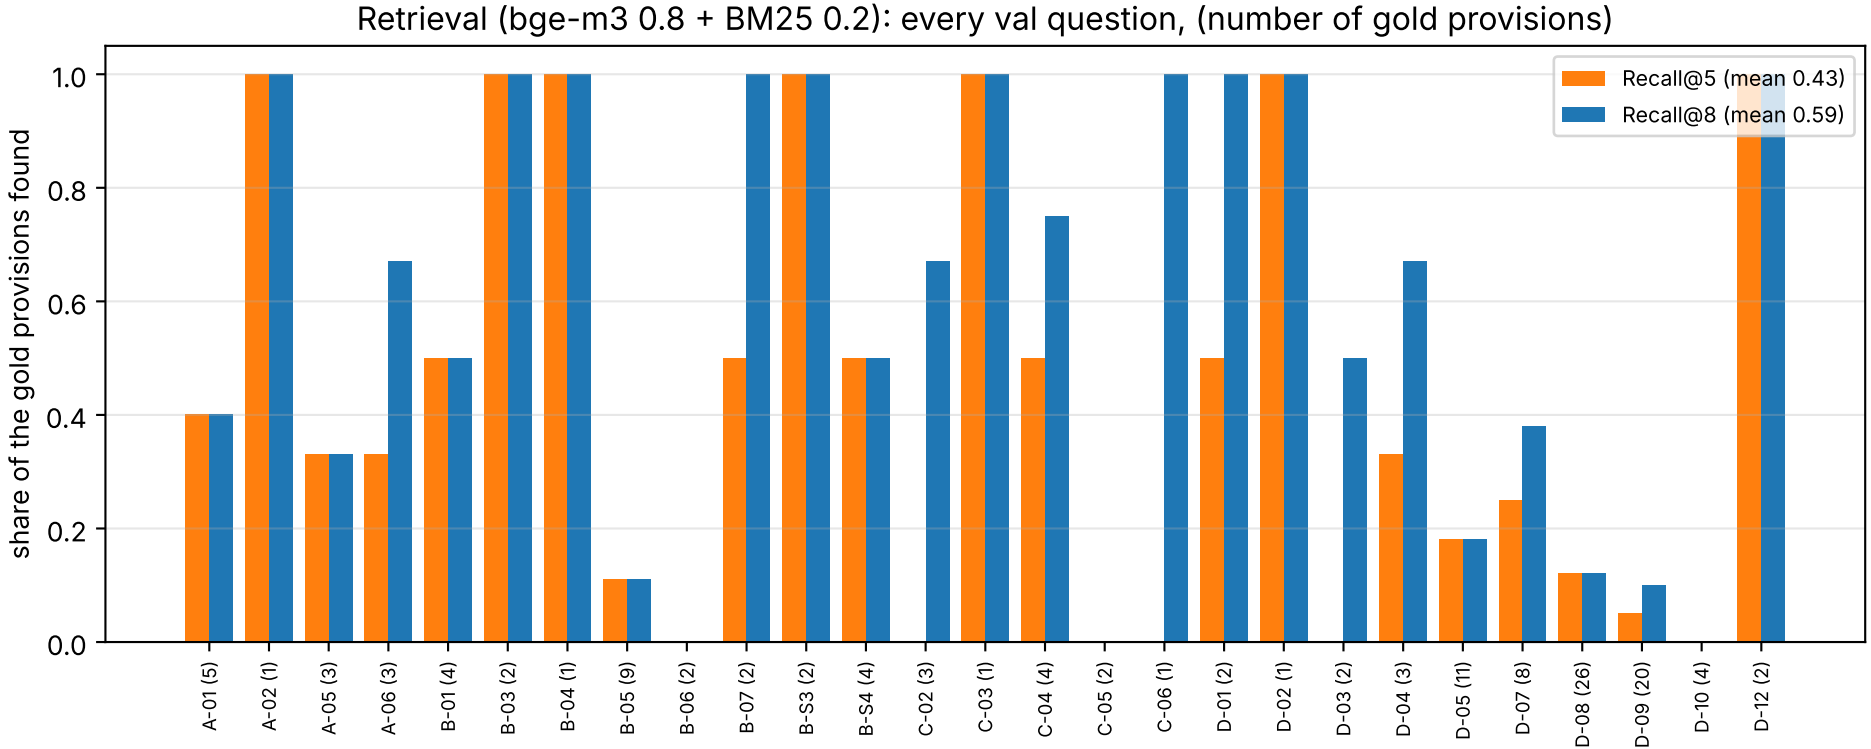
Retrieval (bge-m3 0.8 + BM25 0.2): every val question, (number of gold provisions)
share of the gold provisions found
Recall@5 (mean 0.43)
Recall@8 (mean 0.59)
0.0
0.2
0.4
0.6
0.8
1.0
A-01 (5)
A-02 (1)
A-05 (3)
A-06 (3)
B-01 (4)
B-03 (2)
B-04 (1)
B-05 (9)
B-06 (2)
B-07 (2)
B-S3 (2)
B-S4 (4)
C-02 (3)
C-03 (1)
C-04 (4)
C-05 (2)
C-06 (1)
D-01 (2)
D-02 (1)
D-03 (2)
D-04 (3)
D-05 (11)
D-07 (8)
D-08 (26)
D-09 (20)
D-10 (4)
D-12 (2)

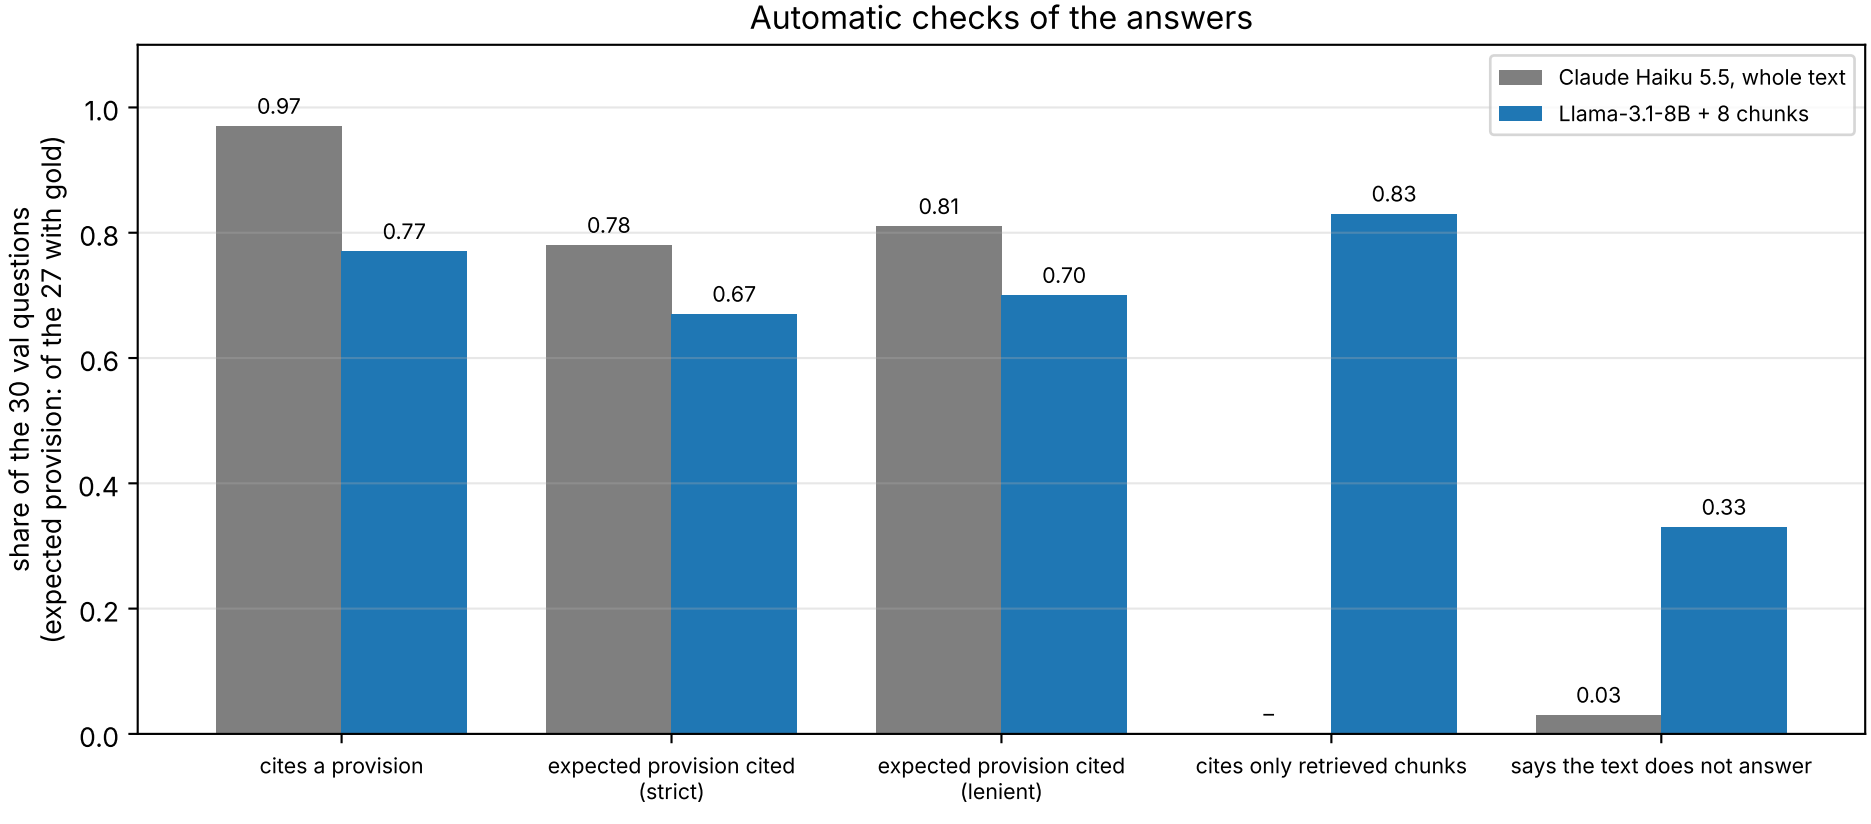
Automatic checks of the answers
share of the 30 val questions
(expected provision: of the 27 with gold)
Claude Haiku 5.5, whole text
Llama-3.1-8B + 8 chunks
1.0
0.8
0.6
0.4
0.2
0.0
0.97
0.77
0.78
0.67
0.81
0.70
–
0.83
0.03
0.33
cites a provision
expected provision cited (strict)
expected provision cited (lenient)
cites only retrieved chunks
says the text does not answer

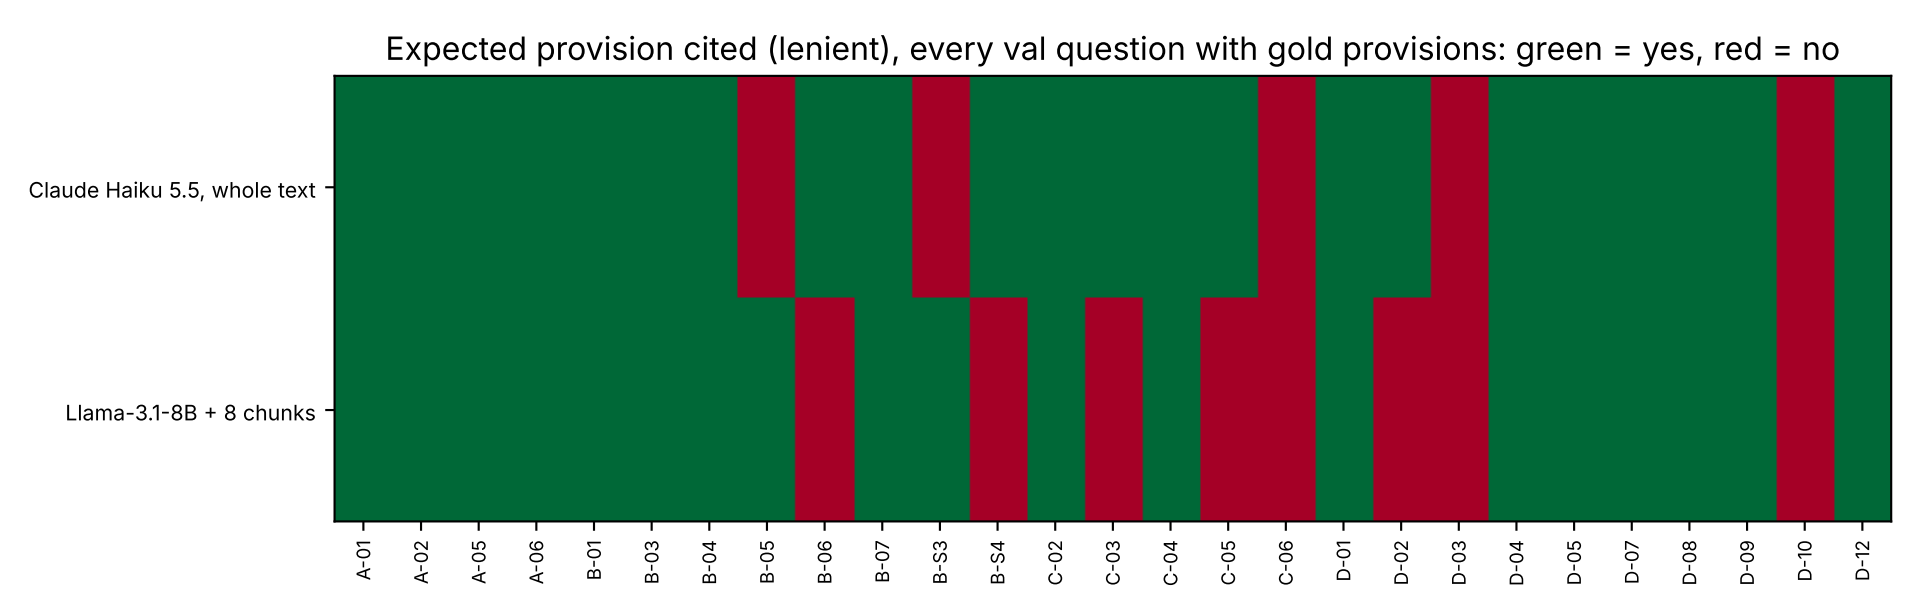
Expected provision cited (lenient), every val question with gold provisions: green = yes, red = no
Claude Haiku 5.5, whole text
Llama-3.1-8B + 8 chunks
A-01
A-02
A-05
A-06
B-01
B-03
B-04
B-05
B-06
B-07
B-S3
B-S4
C-02
C-03
C-04
C-05
C-06
D-01
D-02
D-03
D-04
D-05
D-07
D-08
D-09
D-10
D-12

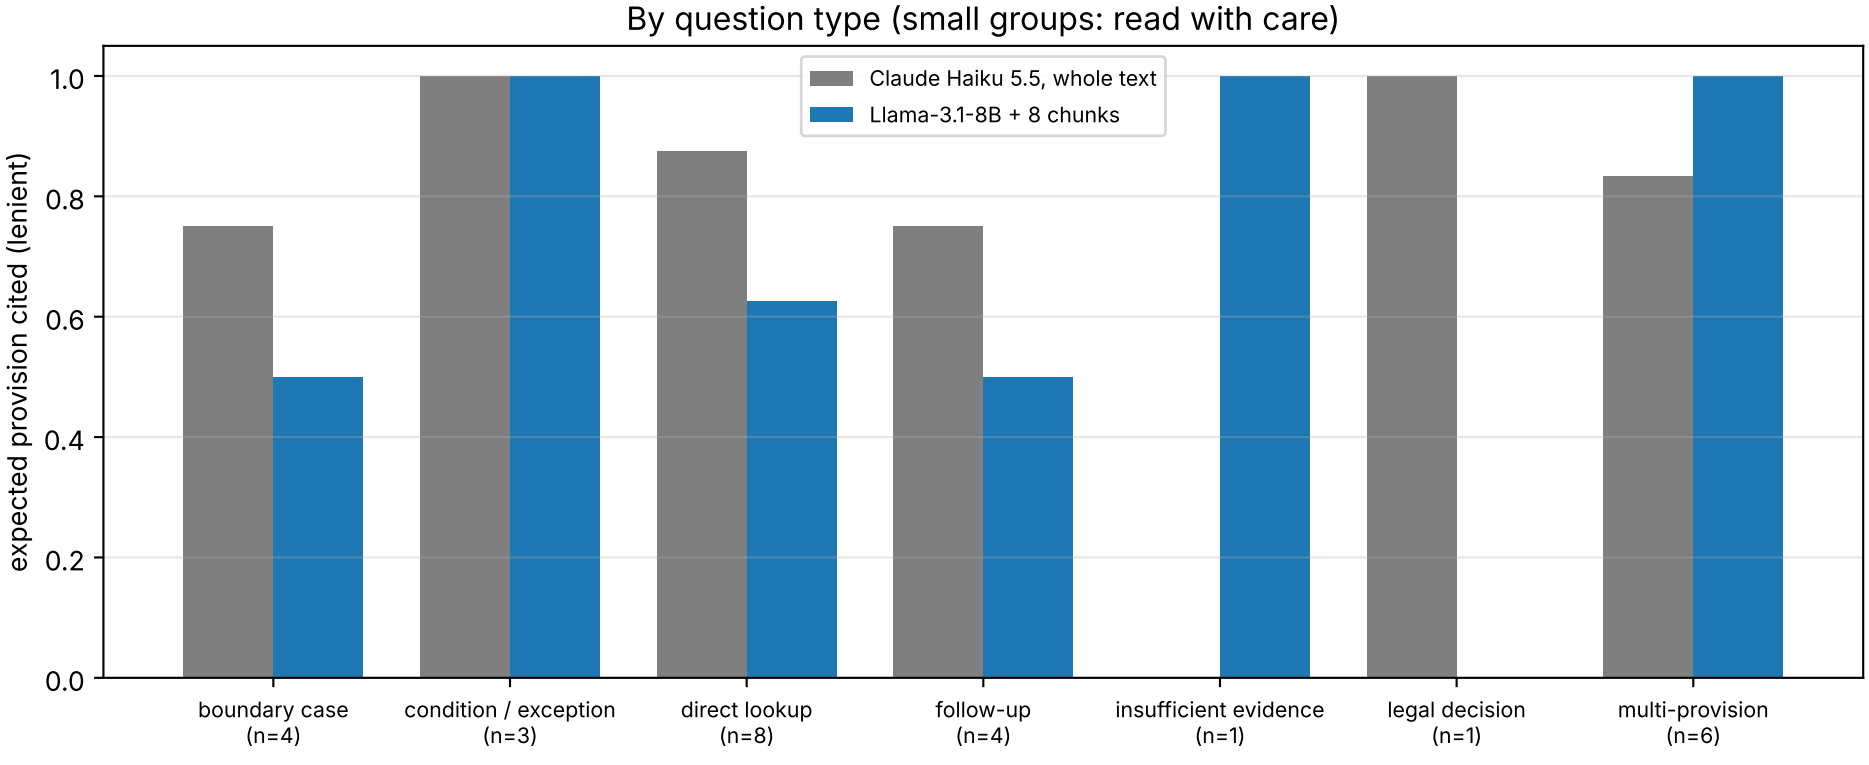
By question type (small groups: read with care)
Claude Haiku 5.5, whole text
Llama-3.1-8B + 8 chunks
expected provision cited (lenient)
1.0
0.8
0.6
0.4
0.2
0.0
boundary case (n=4)
condition / exception (n=3)
direct lookup (n=8)
follow-up (n=4)
insufficient evidence (n=1)
legal decision (n=1)
multi-provision (n=6)

With retrieval: retrieval against answer (27 questions)

	answer cites it	answer does not
gold in the 8 chunks	19	5
no gold in the 8 chunks	0	3